


iMorpheus.ai

Semantic-Preserving BoW Models and Applications

Learning BoW by Y. Li @ iMorpheus.ai

Date Jan. 12, 2018 12:00PM



Journal club介绍与自动驾驶中定位方案相关的论文， 主要关注的方向有：**SLAM**算法、点云数据的处理和压缩、特征地图、传感器数据处理和融合、**GNSS**信号处理等。我们一直关注领域前沿技术，选取得到广泛认可的、或者是在我们的实际使用中结果比较好的论文，与大家分享，共同学习成长。

每周五 北京时间12点

<http://imorpheus.ai/journalclub>



扫码加入无人驾驶技术群



iMorpheus.ai


Background

❖ BoW deals with an image:


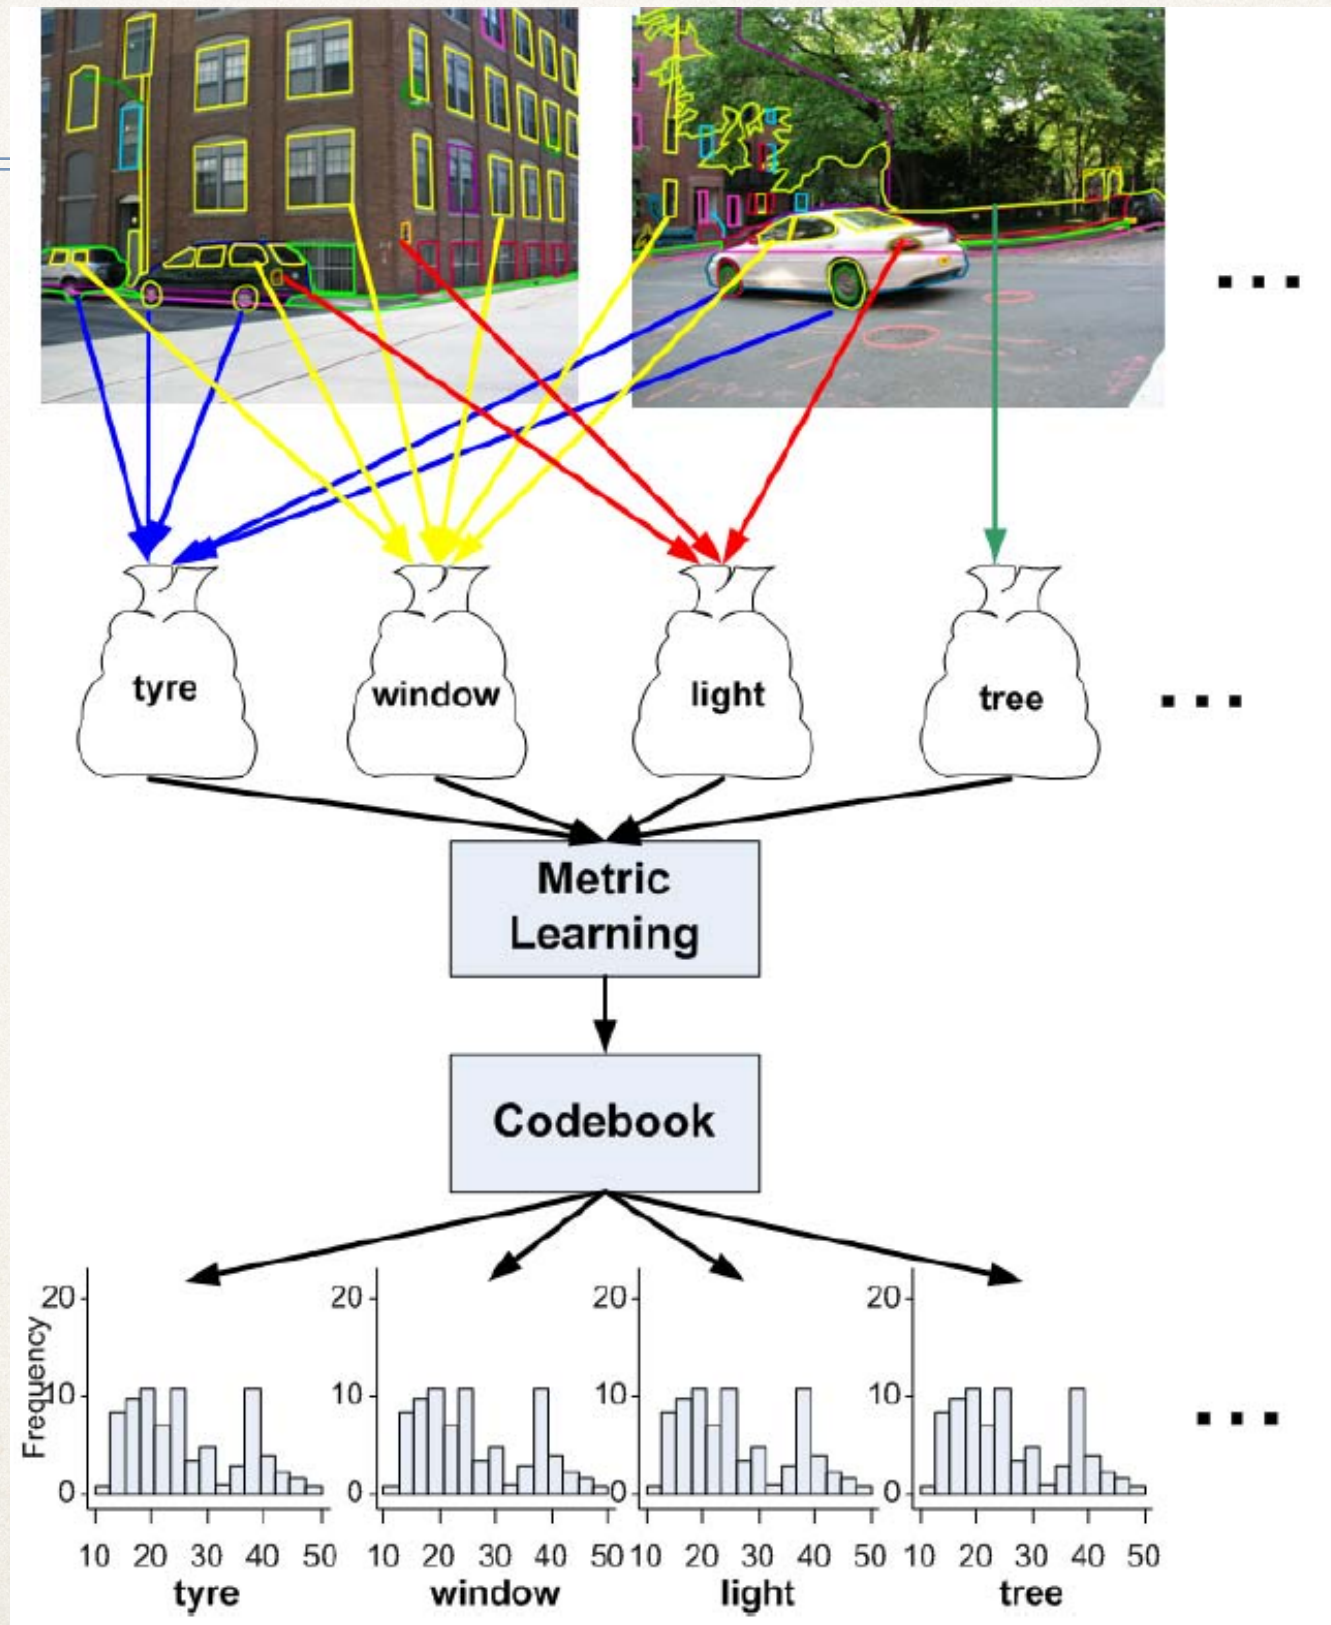
- Interest point detector (e.g. Difference of Gaussians) to detect salient patches / regions in the image.
- Feature descriptor (e.g. SIFT) to represent the local patches / regions as numerical feature vectors.
- To generate a codebook by converting the patches to “codewords”, e.g. applying k-means clustering and defining codewords based on the centers of the clusters.

❖ Drawbacks of BoW:

- Ignorance of spatial information
- Semantics of objects is considerably lost



Semantic Gap ->a distance metric learning method



iMorpheus.ai

Data descriptions

❖ Training data:

- MIT label testbed of images (495 objects, 185 images, 400000 features) , VOC2006


❖ Side information:

- considering pairwise feature instance x_{i1}, x_{i2} ,
- z_{i1} and z_{i2} are binary indicators to indicate whether a feature instance is located at the object region or the background region in the image.
- y_i shows if both x_{i1} and x_{i2} are on the same semantic parts of objects


- ❖ The basic Principle of metric learning:

$$d(x_{i1}, x_{i2}) = \sqrt{(x_{i1} - x_{i2})^\top A (x_{i1} - x_{i2})} \quad (1)$$


- The distances between visual feature vectors of the same semantics should be minimized.
- Distances between feature vectors of different semantics should be maximized.




(a)



(b)



(c)



(d)

Fig. 2. Illustration of side information between objects and feature instances.

I. The distance metric estimation

$$\min_{A \succeq 0, b} \sum_i z_{i1} z_{i2} \xi_i + \frac{\lambda}{2} \text{tr}(AA^\top) \quad (2)$$

$$s.t. \quad y_i(\|x_{i1} - x_{i2}\|_A - b) \leq \xi_i, \xi_i \geq 0, i = 1, \dots, n \quad (3)$$

$$\|A\| = 1/\sqrt{\lambda} \quad (4)$$


- ❖ Mahanalobis distance between two features under metric A
- ❖ Matrix:

$$\underline{X}_1 = [x_{11}, x_{21}, \dots, x_{n1}]^\top \quad \underline{X}_2 = [x_{12}, x_{22}, \dots, x_{n2}]^\top$$

$$\underline{Z}_1 = \text{diag}(z_{11}, z_{21}, \dots, z_{n1}) \quad \underline{Z}_2 = \text{diag}(z_{12}, z_{22}, \dots, z_{n2})$$

$$\underline{Y} = \text{diag}[y_1, \dots, y_n]$$

- ❖ A gradient search algorithm



Algorithm 1 The Semantics-Preserving Metric Learning (SPML) algorithm

INPUT:


- SIFT feature matrix: $X \in \mathbb{R}^{N \times d}$
- pairwise constraint $(x_{i1}, x_{i2}, z_{i1}, z_{i2}, y_i)$, where x_{i1} is the i_1^{th} SIFT feature, z_{i1} indicate whether the location of feature x_{i1} is on the semantic object, and constraints $y_i = \{+1, 0, -1\}$ represents feature x_{i1} and x_{i2} are on the same semantic part of the object, not known, or on different semantic parts.
- regularization parameter λ
- learning rate parameter γ

PROCEDURE:

- 1: initialize metric and threshold: $A = I, b = b_0$
- 2: set iteration step $t = 1$;
- 3: **repeat**
- 4: (1) update the learning rate:
 $\gamma = \gamma/t, t = t + 1$
- 5: (2) update the subset of training instances:
 $\mathcal{S}_t^+ = \{(x_{i1}, x_{i2}, y_i) | (1 + y_i) \|x_{i1} - x_{i2}\|_A^2 > 1\}$
 $\mathcal{S}_t^- = \{(x_{i1}, x_{i2}, y_i) | (1 - y_i) \|x_{i1} - x_{i2}\|_A^2 < 1\}$
 $\mathcal{S}_t = \mathcal{S}_t^+ \cup \mathcal{S}_t^-$
- 6: (3) compute the gradients w.r.t. A
 $\nabla_A \mathcal{L} \leftarrow Z_1 Z_2 (\lambda A + D_X^\top Y^\top D_X),$
 $D_X = X_1 - X_2,$
- 7: (4) compute the gradients w.r.t. b
 $\nabla_b \mathcal{L} \leftarrow \text{tr}(Z_1 Z_2 Y)$
- 8: (5) update metric and threshold:
 $A_{t+1} \leftarrow A_t - \frac{\gamma}{t} \nabla_A \mathcal{L}, \quad b_{t+1} \leftarrow b_t - \frac{\gamma}{t} \nabla_b \mathcal{L}$
- 9: (6) project A back to the PSD cone:
 $A_{t+1} = \sum_{i=1}^d \lambda_i \phi_i \phi_i^\top$
 $A_{t+1} \leftarrow \sum_i \max(0, \lambda_i) \phi_i \phi_i^\top$
- 10: (7) normalize A_{t+1} to satisfy $\|A_{t+1}\| = \frac{1}{\sqrt{\lambda}}$:
 $A_{t+1} \leftarrow \frac{1/\sqrt{\lambda}}{\|A_{t+1}\|} A_{t+1}$
- 11: **until** convergence

OUTPUT:

- feature metric A , threshold variable b



2. Codebook Generation

1. Codebook size assignment ($L_{max}=2500$)

The number of codes \rightarrow visual complexity of an object category \rightarrow the diversity of its associated features

The generative probability of x_j from the object C_i word-bag:

$$p(x_j|C_i) = \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{\|x_j - \hat{x}\|_A^2}{2\sigma^2}}$$

$$\hat{x} = \frac{1}{n_{C_i}} \sum_{x_j \in C_i} x_j$$


Information entropy:

$$H(C_i) = - \sum_{x_j \in C_i} p(x_j|C_i) \log p(x_j|C_i)$$

The number of codes or words:

$$L_{C_i} = \lfloor L_{max} \times \frac{H(C_i)}{\log n_{C_i}} \rfloor$$

2. Codebook generation



Algorithm 2 Codebook Generation Algorithm

INPUT:


- features and their object labels $\{(x, y), x \in \mathcal{X}, y \in \mathcal{C}\}$
- optimized distance metric A
- codebook size assigned for each object $L_{C_i}, i = 1, \dots, M$
- the number of clusters for clustering $K > \max_i L_{C_i}$

PROCEDURE:

- 1: initialize the number of visual words $L = 0$
- 2: **for** $i = 1 : M$ **do**
- 3: clustering features of the i -th object $X_i = \{(x, C) | C = C_i\}$ into K clusters
 $[c_{ij}, r_{ij}] = kmeans(X_i, K)$
- 4: calculate the size of each cluster:
 $S_{ij} = \sum_x \delta(\|x - c_{ij}\|_A, r_{ij})$
- 5: sort clusters by their sizes
 $c_{ij} \leftarrow sort(c_{ij}, S_{ij}) \quad r_{ij} \leftarrow sort(r_{ij}, S_{ij})$
- 6: adopt top L_{C_i} largest clusters as visual words for the category
 $w_{L+j} = c_{ij}, r_{L+j} = r_{ij}, j = 1, \dots, L_{C_i}$
- 7: update the number of visual words $L = L + L_{C_i}$
- 8: **end for**

OUTPUT:

- the centers of visual words w_k and their range radius $r_k, k = 1, \dots, L_{max}$
-



3. Visual word Histogram

❖ Difference from traditional BoW methods:

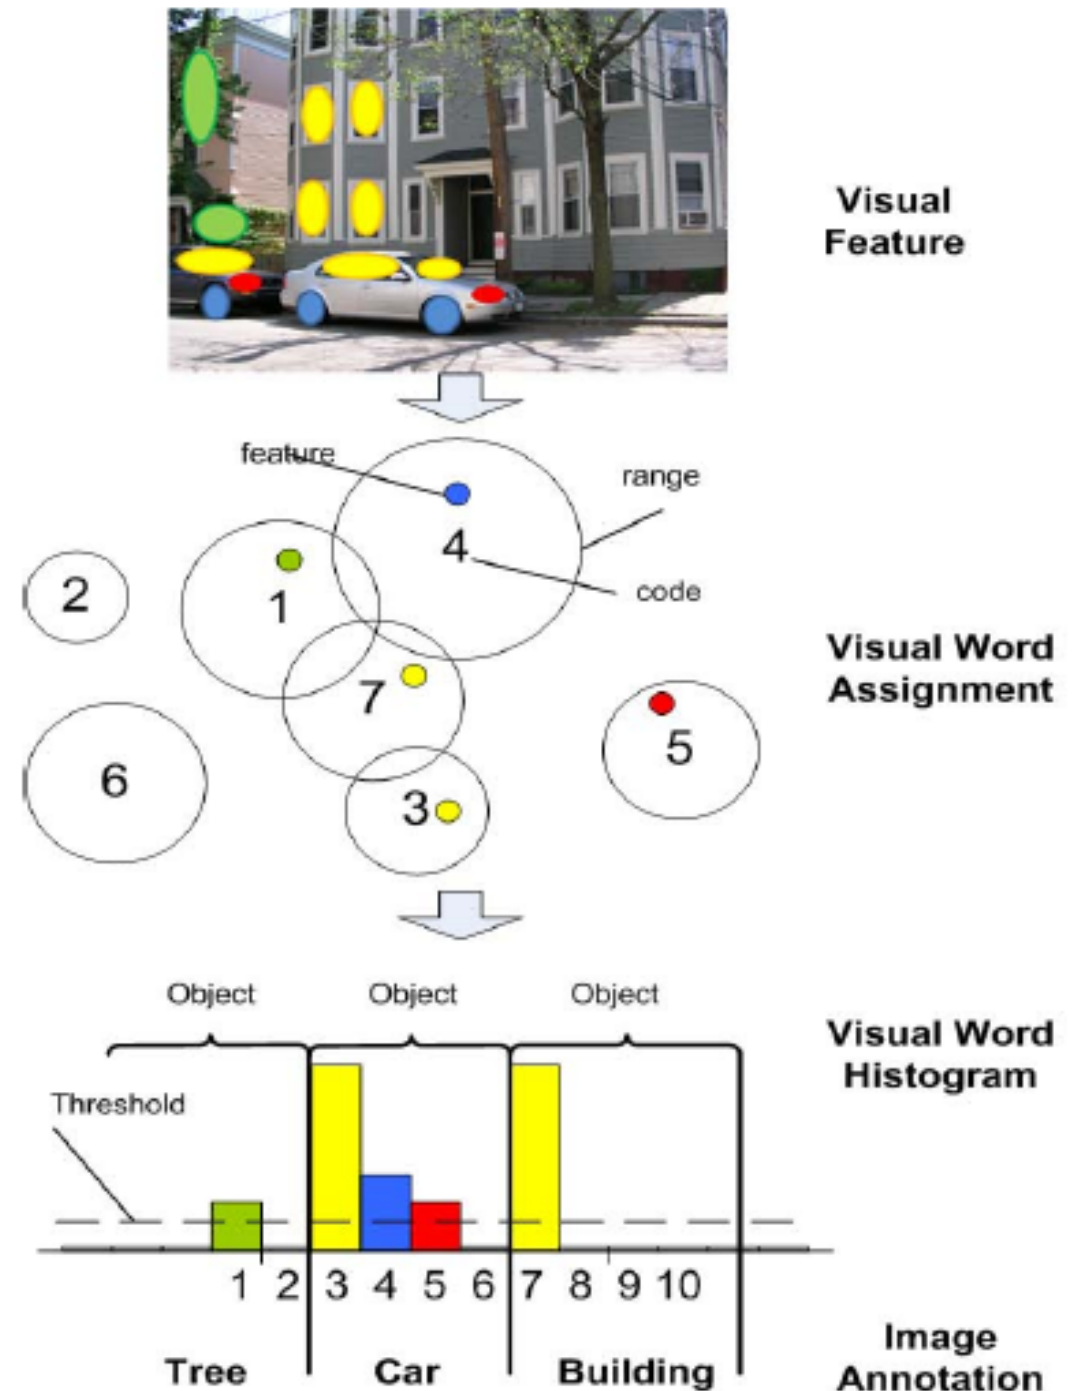
- Traditional BoW assigns a feature to its closest visual words (a cluster).
- In SPBoW, a feature can be assigned to multiple visual words in different objects.

❖ Visual word Histogram:

$$\pi(x, k) = \begin{cases} 1, & \|x - w_k\|_A < r_k; \\ 0, & \text{otherwise.} \end{cases}$$

$$f_I(k) = \sum_{x \in I} \pi(x, k)$$

$$h_I(w_k) = \frac{f_I(k)}{\sum_{v=1}^{L_{max}} f_I(v)}$$



Application (Object Classification)

- ❖ Assuming that we are given a set of labeled image regions. Our goal is to automatically annotate a novel image I

- Generative models


Probability of a visual word w_k appearing in an object C_i , calculated based on the training set:

$$p(w_k|C_i) = \frac{\sum_{\{x|C(x)=C_i\}} \pi(x, k) + 1}{\sum_k \sum_{\{x|C(x)=C_i\}} \pi(x, k) + V}$$

The likelihood of object category C_i appearing in image I can be calculated by a Naive Bayes model :

$$p(C_i|I) \propto p(I|C_i)p(C_i) \propto p(C_i) \prod_k p(w_k|C_i)^{f_I(k)} \quad (12)$$

The top N ranked categories are used to annotate the image



Application (Object Classification)

- ❖ Assuming that we are given a set of labeled image regions. Our goal is to automatically annotate a novel image I .
 - Discriminative models

We are given a set of training images (or image regions) and their semantic categories $\{I_j, C(I_j)\}$, j is 1 to N_{tr} . Each image has visual word histogram $\mathbf{h}_I = [h_I(w_1), h_I(w_2), \dots, h_I(w_{L_{max}})]$, Which is a L_{max} -dimensional vector.

A multi-class classification task, based on the visual word histogram.

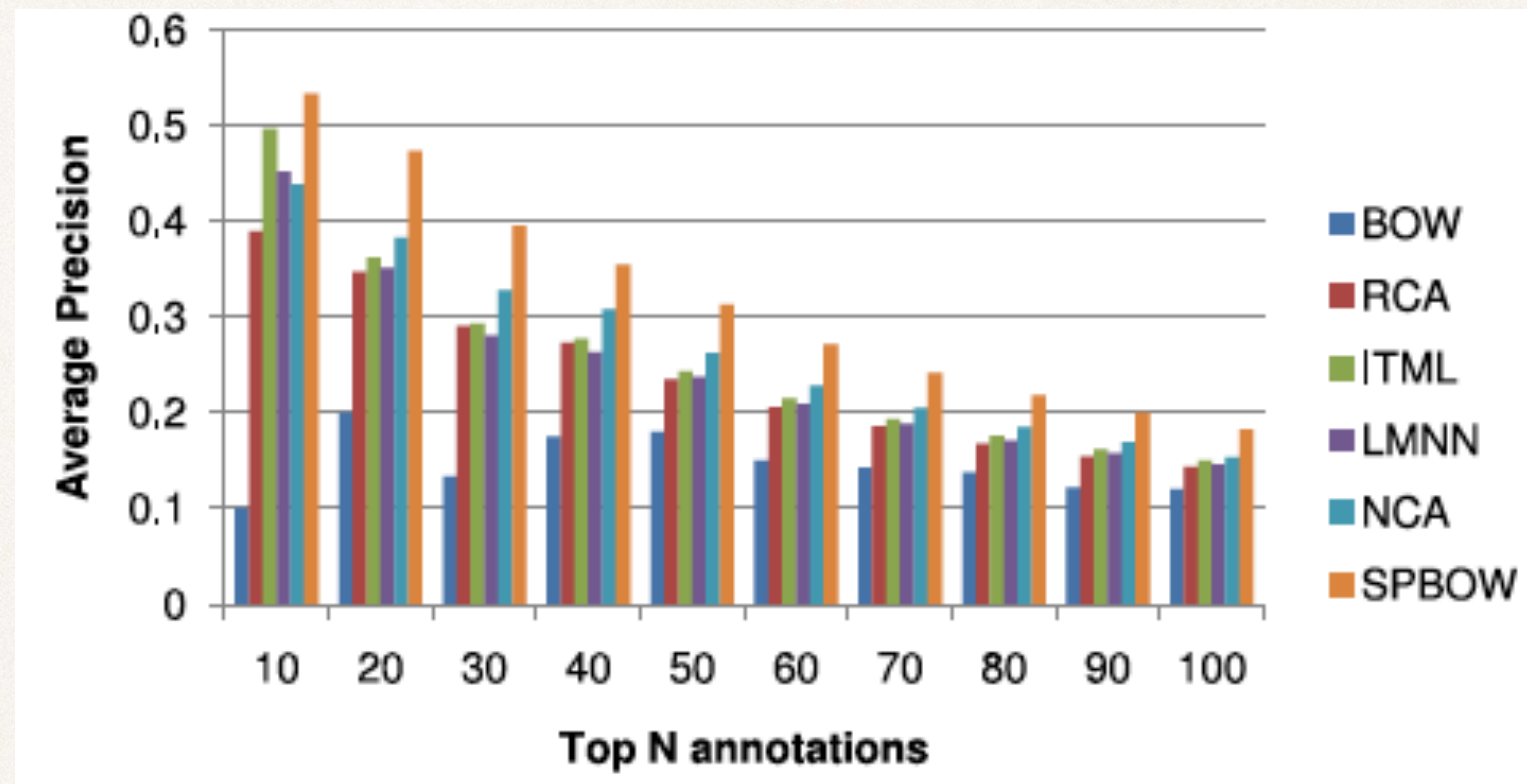
Using a binary SVM classifier to get the weight vector, w and b , for the i th category.

$$\begin{aligned} \min_{\omega, b} \quad & \frac{1}{2} \|\omega\|^2 + C \sum_j \xi_j \\ \text{s.t.} \quad & y_j(i)(\omega \cdot \mathbf{h}_{I_j} - b) \geq 1 - \xi_j, \xi \geq 0, 1 \leq j \leq N_{tr} \end{aligned}$$

A novel test image will be classified by all of the binary SVM classifiers, in which a positive output indicates that a specific object is detected on the image.


Results and Discussion

- ❖ Average Precision
~top N annotation



- ❖ Other influences: Number of constraints, codebook size and object categories (<10, Precision>90%)
- ❖ Codebook Generation:

Method	BoW	RCA	ITML	LMNN	NCA	SPBoW
Time Cost (s)	121	3	96	1759	457	8



❖ **Thank You**

Appendix

True Positive , False Positive , true negative , false negative

Precision其实就是在识别出来的图片中，True positives所占的比率：

$$P = \text{Tp} / (\text{Tp} + \text{Fp})$$

其中的(True positives + False positives)也就是系统一共识别出来多少照片。

在这一例子中，True positives为3，False positives为1，所以Precision值是 $3 / (3 + 1) = 0.75$ 。

意味着在识别出的结果中，飞机的图片占75%。

Recall 是被正确识别出来的飞机个数与测试集中所有飞机的个数的比值：

$$R = \text{Tp} / (\text{Tp} + \text{Fn})$$


Recall的分母是(True positives + False negatives)，这两个值的和，可以理解为一共有多少张飞机的照片。

在这一例子中，True positives为3，False negatives为2，那么Recall值是 $3 / (3 + 2) = 0.6$ 。

意味着在所有的飞机图片中，60%的飞机被正确的识别成飞机。



iMorpheus.ai



iMorpheus.ai Weekly Journal Club

Next Friday, 19/01/2018 12:00PM GMT+8

An introduction to Inertial Navigation

Website : <http://imorpheus.ai>
Email Address : live@imorpheus.ai



iMorpheus.ai